

Monte Carlo Techniques

Alex Zhou

May 2019

1 Introduction

Let $U = (U_1, U_2, \dots)$ represented an infinite sequence of coin tosses with $U_i = 1$ if the i th toss is heads and $U_i = 0$ if it is tails. Suppose the $(U_i)_i$ are independent and identically distributed (iid), and the probability of heads is $p \in (0, 1)$ so the probability of tails is $q = 1 - p$. Define a real-valued random variable $X = f(U)$ taking values in $[0, 1]$ by

$$f(U) = \sum_{i=1}^{\infty} \frac{U_i}{2^i},$$

which can be thought of as a binary expansion of X , (although note that rational numbers do not have a unique binary expansion). Define the cumulative distribution function

$$F(x) = \Pr(X \leq x).$$

We intend to estimate F and reveal some interesting properties.

2 Monte Carlo Simulation of the Empirical CDF

Fix $n \in \mathbf{N}$. Generate a finite sequence $U^n = (U_1, \dots, U_n)$ and compute $X_n = \sum_{i=1}^n U_i/2^i$. Repeat this N times to generate a random sample X_1^n, \dots, X_N^n which we can use to define the empirical cumulative distribution function

$$\hat{F}(x) = \frac{1}{N} \sum_{j=1}^n \mathbf{1}[X_j^n \leq x],$$

where $\mathbf{1}[A]$ is the indicator function for a set A . This should be an approximation for the actual cumulative distribution function $F(x)$. By the strong law of large numbers, as $N \rightarrow \infty$, the empirical distribution $\hat{F}(x)$ converges almost surely to the true cumulative distribution function $F(x)$, and this convergence has rate $1/\sqrt{N}$. Taking $N = 10000$ ensures that the number of samples is sufficiently large so that our error is not noticable on the plot. The following Python code implements this procedure:

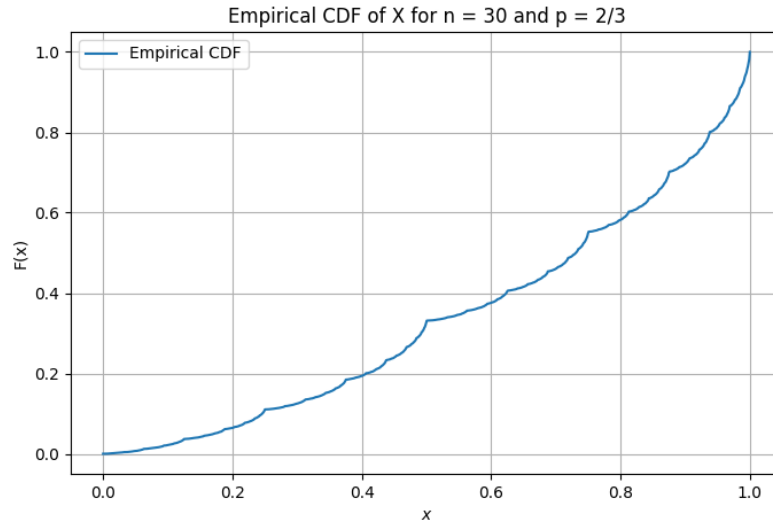


Figure 1: $N = 10000$

```

1  def generate_sample(n, p, N):
2      # Generate an N by n matrix of Bernoulli trials
3      U = np.random.binomial(1, p, (N, n))
4      powers_of_two = 2 ** np.arange(1, n + 1)
5      # Returns a weighted sum of each row
6      X = U @ (1 / powers_of_two)
7      return X
8
9  def F_hat(sample, x_values):
10     N = len(sample)
11     return [np.sum(sample <= x) / N for x in x_values]
12
13  def ecdf(n, p, N):
14     sample = generate_sample(n, p, N)
15     x_values = np.linspace(0, 1, 2**n + 1)
16     y_values = F_hat(sample, x_values)
17     return x_values, y_values

```

Using the generated sample we can plot the empirical cumulative distribution in the case where $n = 30$ and $p = \frac{2}{3}$.

3 Exact Values of the CDF

For some values of x , the cumulative distribution function $F(x) = \Pr(X \leq x)$ can be explicitly computed. In particular, let

$$x = \sum_{i=1}^n \frac{x_i}{2^i},$$

for fixed $n \in \mathbf{N}$ and a sequence $x_1, \dots, x_n \in \{0, 1\}$. That is, x has a finite binary expansion. We need to find the probability that the random infinite binary expansion of X is less than or equal to the finite binary expansion of x . This can be done by comparing digit by digit which we do recursively. Let us write the binary representations $x = [0 \cdot x_1 \dots x_n]_2$ and $X = [0 \cdot U_1 \dots]_2$.

We first note that x admits two binary representations if $x > 0$ and $x_n = 1$, namely $x = [0 \cdot x_1 \dots x_{n-1}1000 \dots]_2$ and $x = [0 \cdot x_1 \dots x_n 0111 \dots]_2$, (the second expansion is not valid for $x_n = 0$). Let $s = [0s_1, \dots s_n 000 \dots]_2$ be one binary expansion for x . The event $X = x$ via this sequence means $U_i = s_i$ for all i . The probability is then

$$\Pr(X = x) = \prod_{i=1}^{\infty} \Pr(U_i = s_i) = \lim_{N \rightarrow \infty} (1-p)^{N-n} \prod_{i=1}^n \Pr(U_i = s_i) = 0,$$

and the same equality holds for the other binary expansion. Therefore $F(x) = \Pr(X < x)$.

For $x = [0.x_1 \dots x_n]_2$, consider the first coin toss U_1 :

1. If $U_1 < x_1$, then $U_1 = 0$ and $x_1 = 0$ so $X \leq 1/2$ whilst $x \geq 1/2$, so $X < x$ is always true and we have $\Pr(X < x | U_1 < x_1) = (1-p)x_1$.
2. If $U_1 > x_1$, then $X < x$ is always false hence, $\Pr(X < x | U_1 > x_1) = 0$.
3. If $U_1 = x_1$, then the condition $X < x$ now fall to the next digits U_2 and x_2 . We have the Bernoulli mass function $\Pr(U_1 = x_1) = p^{x_1}(1-p)^{1-x_1}$, yielding the recursion

$$\Pr(X < x | U_1 = x_1) = p^{x_1}(1-p)^{1-x_1} F([0 \cdot 0x_2 \dots x_n]_2).$$

Altogether, summing the conditional probabilities we obtain

$$F(x) = x_1(1-p) + p^{x_1}(1-p)^{1-x_1} F([0 \cdot 0x_2 \dots x_n]_2).$$

Let $x^{(k)} = [0 \cdot 0x_k \dots x_n]_2$ and $P_j = p^{x_j}(1-p)^{1-x_j}$. We solve the recurrence

$$F(x^{(k)}) = x_k(1-p) + P_k F(x^{(k+1)})$$

by backwards iteration on k . The base case $k = n+1$ is trivial $F(0) = F(X = 0) = 0$, so

$$F(x^{(n)}) = x_n(1-p) + P_n \cdot 0 = x_n(1-p),$$

and continuing the backwards substitution yield

$$F(x^{(1)}) = x_1(1-p) + x_2(1-p)P_1 + \dots + x_n(1-p)P_1 \dots P_n = \sum_{k=1}^n x_k(1-p) \prod_{j=1}^{k-1} P_j,$$

where the empty product is 1. Hence, we have

$$F(x) = \sum_{k=1}^n x_k(1-p) \prod_{j=1}^{k-1} p^{x_j} (1-p)^{1-x_j},$$

is the desired formula for the cumulative distribution function. We implement this algorithm below in a similar fashion to the empirical distribution function:

```

1  def binary_sequences(n):
2  return np.array([
3      [(i >> (n - j - 1)) & 1 for j in range(n)]
4  for i in range(2**n)])
5
6  def F(p, sequence):
7      F_x = 0
8      for k in range(len(sequence)):
9          if sequence[k] == 1:
10             product = 1
11             for j in range(k):
12                 if sequence[j] == 1:
13                     product *= p
14                 else:
15                     product *= (1 - p)
16             F_x += (1 - p) * product
17 return F_x
18
19 def cdf(n, p):
20     sequences = binary_sequences(n)
21     x_values = np.linspace(0, 1, 2**n + 1)
22     y_values = np.array([F(p, seq) for seq in sequences] + [1.0])
23     return x_values, y_values

```

For the empirical distribution, the time-complexity is $O(Nn)$ provided the number of x values we test is dominated by the sample size N . This is because generating the sample runs in $O(Nn)$ time as it generates N sequences of length n , whilst computing \hat{F} for N samples at M test points has time $O(NM)$. This could potentially be improved by sorting the sample once $O(N \log N)$, then using binary search to compute at x the value of the CDF $O(M \log N)$.

```

1  def ecdf(n, p, N):
2      sample = generate_sample(n, p, N)

```

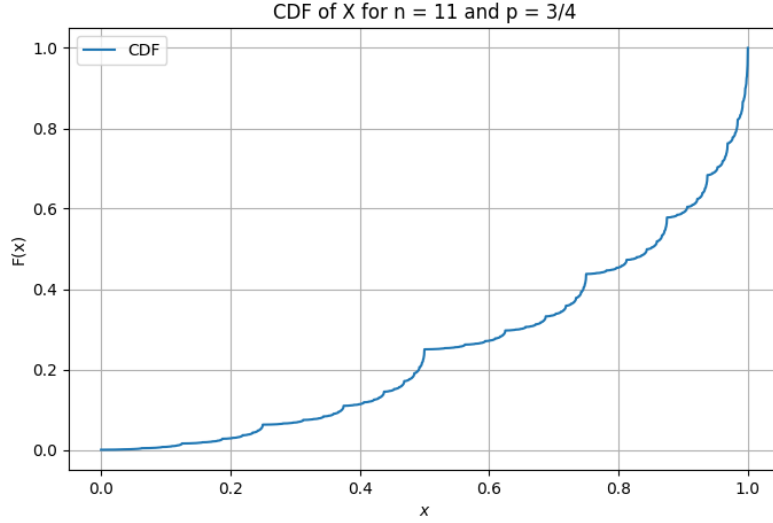


Figure 2: True CDF for X , estimated at the dyadic rationals.

```

3 sorted_sample = np.sort(sample)
4 x_values = np.linspace(0, 1, 2**n + 1)
5 y_values = np.searchsorted(sorted_sample, x_values, side='right') / N
6 return x_values, y_values

```

However, the time-complexity of our cumulative distribution function algorithm is $O(n^2 2^n)$. Generating the binary sequences requires $O(n 2^n)$ time, as we loop over 2^n integers and perform n bitwise operations. Computing F in the worst case is $O(n^2)$, as the double loop scales alike the triangle numbers, and it is called 2^n times. We can make an improvement when computing F since the product can be updated iteratively for each kF , reducing time complexity to $O(n)$.

```

1 def F(p, sequence):
2     F_x = 0
3     product = 1
4     for k in range(len(sequence)):
5         if sequence[k] == 1:
6             F_x += (1 - p) * product
7         if sequence[k] == 1:
8             product *= p
9         else:
10            product *= 1 - p
11 return F_x

```

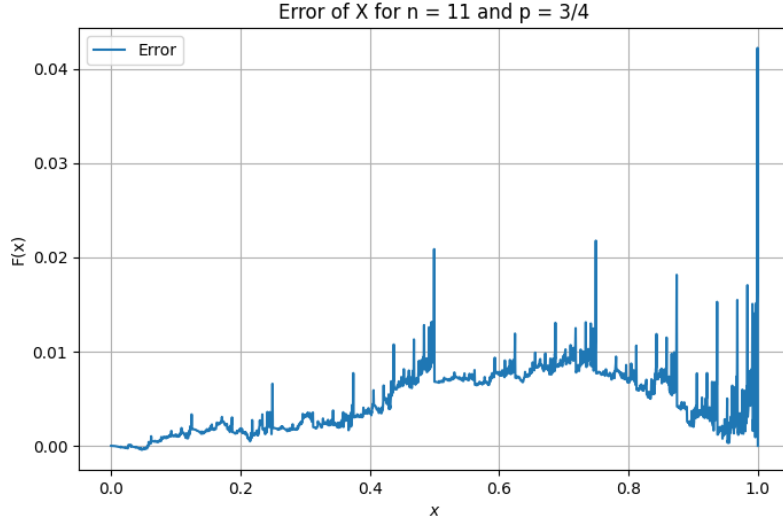


Figure 3: The difference between the ECDF and CDF for X .

Nevertheless, for large n , this algorithm becomes infeasible to use.

Indeed, $n = 11$ is quite small, but $2^{11} = 2048$ is relatively computationally expensive.

4 Properties of the CDF

Recall that F is continuous at c if

$$\lim_{x \rightarrow c^-} F(x) = F(c) = \lim_{x \rightarrow c^+} F(x).$$

Note that the right equality always holds for a valid CDF, so continuity is equivalent to the left equality. Since $F(c) - \lim_{x \rightarrow c^-} F(x) = \Pr(X = c)$, we have that $F(x)$ is continuous at c if and only if $\Pr(X = c) = 0$.

Let c have a finite binary expansion. We first claim that $F(x) = \Pr(X \leq x)$ is continuous at $x = c$. Write $X = [0 \cdot U_1 U_2 \dots]_2$ and

$$c = \frac{k}{2^n} = [0 \cdot c_1 \dots c_n]_2.$$

1. If $c = 0$, then

$$\Pr(X = 0) = \Pr\left(\bigcap_{i=1}^{\infty} \{U_i = 0\}\right) = \prod_{i=1}^{\infty} \Pr(U_i = 0) = \prod_{i=1}^{\infty} (1 - p).$$

Since $(1 - p) \in (0, 1)$, then $\lim_{N \rightarrow \infty} (1 - p)^N = 0$ so $\Pr(X = 0) = 0$.

2. If $c = 1$, then by symmetry, $\lim_{N \rightarrow \infty} p^N = 0$, so $\Pr(X = 1) = 0$.
3. If $c \in (0, 1)$, then

$$\Pr(X = c) = \Pr\left(\bigcap_{i=1}^{\infty} \{U_i = c_i\}\right) = \prod_{i=1}^n \Pr(U_i = c_i) \prod_{i=n+1}^{\infty} \Pr(U_i = 0).$$

The first term is finite, and the second term converges to 0 from above. Similarly, by considering the other binary representation of c we obtain $\Pr(X = c) = 0$.

From our previous plots, they seem to suggest that the cumulative distribution function is also continuous at all other points. Such points have infinite binary expansion $c = [0 \cdot c_1 c_2 \dots]_2$, which must contain an infinite number of zeros or ones. Indeed, if the expansion is eventually constant, then it must be a dyadic rational number. Now,

$$\Pr(X = c) = \Pr\left(\bigcap_{i=1}^{\infty} \{U_i = c_i\}\right) = \prod_{i=1}^{\infty} \Pr(U_i = c_i).$$

Define the partial product $P_N = \prod_{i=1}^N \Pr(U_i = c_i)$. Each term $\Pr(U_i = c_i)$ is either $1 - p$ when $c_i = 0$ or p when $c_i = 1$. Let $r = \max\{p, q\}$. Then

$$0 \leq P_n = \prod_{i=1}^N \Pr(U_i = c_i) \leq r^N \in (0, 1),$$

so $\lim_{N \rightarrow \infty} r^N = 0$ hence $\lim_{N \rightarrow \infty} P_N = 0$ by the sandwich theorem. Therefore, we obtain $\Pr(X = c) = 0$ as required. This means that the CDF $F(x)$ is continuous on $[0, 1]$.

Recall that F is left-differentiable at c if

$$\lim_{\delta \rightarrow 0^-} \frac{F(c + \delta) - F(c)}{\delta}$$

exists and is finite. Likewise, F is right-differentiable at c if

$$\lim_{\delta \rightarrow 0^+} \frac{F(c + \delta) - F(c)}{\delta}.$$

We say that F is differentiable at c if it is both left and right differentiable at c . To estimate the difference quotients for dyadic rationals δ in a neighbourhood of c , we shall use the following code:

```

1 def difference_quotients(p, c, n):
2     delta = 1 / (2 ** n)
3     m_values = np.arange(-100, 100 + 1)
4     m_values = m_values[m_values != 0]
5 
```




Figure 4: Difference quotient in a neighbourhood of $c = 9/16$.

```

6     x_values = c + m_values * delta
7     x_values = x_values[(x_values >= 0) & (x_values <= 1)]
8
9     F_c = F(p, binary_expansion(c, n))
10    F_x_values = np.array([F(p, binary_expansion(x, n)) for x in x_values])
11    dq_values = (F_x_values - F_c) / (m_values * delta)
12
13    return x_values, dq_values

```

Set $c = 9/16$ and $p = 3/4$. We shall take $\delta = 1/2^{11}$ as before to obtain a decent approximation of the derivative.

From the plot, if $p > 1/2$ then it seems that F is not left-differentiable at c , since the difference quotient blows-up, whilst it is right-differentiable at c with vanishing derivative. The opposite happens for $p < 1/2$. We conjecture the following result: At a dyadic rational point c , F is left-differentiable at c if $0 < p \leq 1/2$ and right-differentiable at c if $1/2 \leq p < 1$.

Let $\delta_L = \frac{1}{2^L}$ and write $c = [0 \cdot c_1 \dots c_n]_2 \in [0, 1]$ for $L > n$.

- Right-differentiability of $c \in [0, 1]$: Consider the first expansion given by $c = [0 \cdot c_1 \dots c_{n-1} 1000 \dots]_2$. The interval $[c, c + \delta_L]$ consists of numbers whose binary expansions agree with that of c in the first n positions and are equal to 0 from position $n + 1$ up to position L . Therefore

$$\Pr(X \in [c, c + \delta_L]) = \prod_{i=1}^n \Pr(U_i = c_i) \prod_{i=n+1}^L \Pr(U_i = 0) = C_n (1-p)^{L-n},$$

where $C_n := \prod_{i=1}^n \Pr(U_i = c_i)$. The difference quotient is now

$$\frac{C_n (1-p)^{L-n}}{\delta_L} = C_n (1-p)^{-n} (2(1-p))^L.$$

Taking $L \rightarrow \infty$, we see that this limit is finite if and only if $p \leq 1/2$.

- Left-differentiability of $c \in (0, 1]$: Consider the second expansion given by $c = [0 \cdot c_1 \dots c_{n-1} 0111 \dots]_2$. The interval $[c - \delta_L, c]$ consists of numbers whose binary expansions agree with that of c in the first n positions and are equal to 1 from position $n + 1$ up to position L . Therefore

$$\Pr(X \in [c - \delta_L, c]) = \prod_{i=1}^n \Pr(U_i = c_i) \prod_{i=n+1}^L \Pr(U_i = 0) = C_n p^{L-n},$$

where $C_n := \prod_{i=1}^n \Pr(U_i = c_i)$. The difference quotient is now

$$\frac{C_n p^{L-n}}{\delta_L} = C_n p^{-n} (2p)^L.$$

Taking $L \rightarrow \infty$, we see that this limit is finite if and only if $p \geq 1/2$.

To summarise, F is right-differentiable at a dyadic rational $c \in [0, 1)$ if and only if $p \in (0, 1/2]$ with right-derivative 1 if $p = 1/2$ and 0 otherwise. Likewise F is left-differentiable at a dyadic rational $c \in (0, 1]$ if and only if $p \in [1/2, 1)$ with left-derivative 1 if $p = 1/2$ and 0 otherwise. This also means that F is differentiable at c if and only if $p = 1/2$.

5 Monte Carlo Methods in Python

NumPy in particular is very useful for Monte Carlo methods, owing to its probability distribution functions on arrays.